

Azure Monitoring and Scaling Practice Lab

1. Azure Monitoring

Azure Monitoring is all about **observing and analyzing** the health, performance, and availability of your applications, infrastructure, and resources in the Azure cloud.

- **Tool used:** Azure Monitor is the main service.
- **What it does:**
 - Collects **metrics** (e.g., CPU %, memory, request rates, latency).
 - Collects **logs** (e.g., errors, exceptions, activity logs, security logs).
 - Provides **alerts** (e.g., notify if CPU > 80% for 5 minutes).
 - **Dashboards & Insights** – helps visualize data in real-time.
 - Can integrate with **Application Insights** (for app-level monitoring) and **Log Analytics** (for deep querying).

👉 Example: If your Web API in Azure App Service suddenly slows down, Azure Monitor can alert you that **response time > 3 seconds**, helping you fix it before users complain.

2. Azure Scaling

Scaling in Azure means adjusting the **resources** (like CPU, memory, or instances of an app) to meet demand.

This ensures **performance** during high traffic and **cost savings** during low usage.

There are two main types:

a) Vertical Scaling (Scale Up/Down)

- Increasing or decreasing the **power of the resource**.
- Example: Changing your App Service Plan from **Basic (1 vCPU, 1.75 GB RAM)** to **Premium (4 vCPU, 14 GB RAM)**.

b) Horizontal Scaling (Scale Out/In)

- Increasing or decreasing the **number of resource instances**.
- Example: Running your API on **2 instances during low traffic** and automatically scaling to **10 instances during peak hours**.
- Often managed by **Azure Autoscale** rules (e.g., scale out when CPU > 70%).

👉 Scaling is what keeps apps **responsive** and **cost-effective**.

✅ **In short:**

- **Azure Monitoring** = "Keep an eye on everything, get alerts, analyze issues."
 - **Azure Scaling** = "Automatically adjust resources to match demand."
-

Let's build a **hands-on lab exercise** where you'll practice both **Azure Monitoring** and **Auto Scaling** using an **Azure App Service** (since it's the simplest way to start).

Lab Exercise – Azure Monitoring & Scaling with App Service

Prerequisites

- An **Azure subscription**
 - A **sample .NET Web API or Angular app** (or any app you can deploy to Azure App Service)
 - Visual Studio / VS Code or Azure CLI
-

Step 1: Create a Resource Group

1. Go to **Azure Portal** → Search for **Resource Groups**.
 2. Click **+ Create**.
 3. Enter:
 - Resource Group name: `AppServiceLabRG`
 - Region: Select nearest to you (e.g., Central India).
 4. Click **Review + Create** → **Create**.
-

Step 2: Create an App Service

1. In Azure Portal, search for **App Services** → **+ Create**.
2. Fill details:
 - Subscription: Your subscription
 - Resource Group: `AppServiceLabRG`
 - Name: `myappservice-lab` (must be unique)
 - Publish: **Code**
 - Runtime: **.NET 8 (LTS)** or **Node.js/Python** if you prefer
 - OS: **Windows**
 - Region: Same as Resource Group
3. Select **App Service Plan**:
 - Click **Create new** → Choose **B1 (Basic)** (allows scaling).

4. Click **Review + Create** → **Create**.
-

Step 3: Deploy Your Application

- From **Visual Studio**:
 - Right-click project → **Publish** → Select **Azure** → **App Service (Windows)**.
- Or use **Azure CLI**:
- ```
az webapp up --name myappservice-lab --resource-group AppServiceLabRG --runtime "DOTNET:8"
```

Once deployed, open <https://myappservice-lab.azurewebsites.net> in browser.

---

## Step 4: Enable Monitoring

1. Open your App Service → Left menu → **Application Insights**.
  2. Click **Turn On Application Insights**.
  3. Select **Create New** and link to your app.
  4. Choose **Log Analytics Workspace** → Create a new one.
  5. Once enabled, you'll get:
    - **Live Metrics** (requests/sec, response time, failures).
    - **Availability Tests** (ping app every 5 mins).
    - **Alerts** (CPU > 80%, app errors, etc.).
- 

## Step 5: Create an Alert Rule

1. Go to App Service → **Alerts** → + **Create Alert Rule**.
2. Select **Scope**: Your App Service.
3. Select **Condition**:
  - **Signal**: Percentage CPU > 70
  - **Operator**: Greater than
  - **Aggregation**: Average over 5 minutes
4. Select **Action Group**:
  - Create new → Add your **Email ID**.
5. Click **Review + Create**.

✓ Now you'll get an **email alert** when CPU > 70%.

---

## Step 6: Configure Auto Scaling

1. Go to App Service → **Scale out (App Service Plan)**.

2. Select **Custom Autoscale**.
  3. Add Rule:
    - Condition: CPU % > 70 for 5 minutes → **Increase instance count by 1**.
    - Condition: CPU % < 30 for 10 minutes → **Decrease instance count by 1**.
  4. Set **Min Instances = 1** and **Max Instances = 5**.
  5. Save.
- 

## Step 7: Test Auto Scaling

- Use a load testing tool (like **Azure Load Testing**, JMeter, or ab command). Example:
  - `ab -n 10000 -c 200 https://myappservice-lab.azurewebsites.net/`
  - Watch metrics in **Application Insights** → **Live Metrics**.
  - Once CPU > 70%, Azure will automatically add more instances.
- 

## Step 8: Review Logs & Metrics

- Go to **Log Analytics** → **Logs** and run a query:
  - `requests`
  - `| summarize count() by bin(timestamp, 1m), resultCode`
  - View how many requests per minute were served.
  - Check **Alerts** to confirm scaling & email notifications.
- 



## End Result

- You deployed a Web App.
- Enabled **Application Insights Monitoring**.
- Set up **CPU Alerts**.
- Configured **Auto Scaling** rules.
- Tested with **load** to trigger scaling.